

# assignment4

May 28, 2024

## 1 Assignment 4: Self-Attention for Vision

For this assignment, we're going to implement self-attention blocks in a convolutional neural network for CIFAR-10 Classification.

## 2 Part I. Preparation

First, we load the CIFAR-10 dataset. This might take a couple minutes the first time you do it, but the files should stay cached after that.

```
[41]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torch.utils.data import sampler

import torchvision.datasets as dset
import torchvision.transforms as T

import numpy as np
```

```
[42]: NUM_TRAIN = 49000

# The torchvision.transforms package provides tools for preprocessing data
# and for performing data augmentation; here we set up a transform to
# preprocess the data by subtracting the mean RGB value and dividing by the
# standard deviation of each RGB value; we've hardcoded the mean and std.
transform = T.Compose([
    T.ToTensor(),
    T.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010))
])

# We set up a Dataset object for each split (train / val / test); Datasets load
# training examples one at a time, so we wrap each Dataset in a DataLoader which
# iterates through the Dataset and forms minibatches. We divide the CIFAR-10
# training set into train and val sets by passing a Sampler object to the
# DataLoader telling how it should sample from the underlying Dataset.
cifar10_train = dset.CIFAR10('./data/datasets', train=True, download=True,
```

```

        transform=transform)
loader_train = DataLoader(cifar10_train, batch_size=64,
                           sampler=sampler.SubsetRandomSampler(range(NUM_TRAIN)))

cifar10_val = dset.CIFAR10('./data/datasets', train=True, download=True,
                           transform=transform)
loader_val = DataLoader(cifar10_val, batch_size=64,
                        sampler=sampler.SubsetRandomSampler(range(NUM_TRAIN,
↪50000))))

cifar10_test = dset.CIFAR10('./data/datasets', train=False, download=True,
                             transform=transform)
loader_test = DataLoader(cifar10_test, batch_size=64)

```

Files already downloaded and verified  
Files already downloaded and verified  
Files already downloaded and verified

You have an option to **use GPU by setting the flag to True below**. It is not necessary to use GPU for this assignment. Note that if your computer does not have CUDA enabled, `torch.cuda.is_available()` will return False and this notebook will fallback to CPU mode.

The global variables `dtype` and `device` will control the data types throughout this assignment.

```

[43]: USE_GPU = True

dtype = torch.float32 # we will be using float throughout this tutorial

if USE_GPU and torch.cuda.is_available():
    device = torch.device('cuda')
else:
    device = torch.device('cpu')

# Constant to control how frequently we print train loss
print_every = 100

print('using device:', device)

```

using device: cuda

## 2.1 Flatten Function

```

[44]: def flatten(x):
        N = x.shape[0] # read in N, C, H, W
        return x.view(N, -1) # "flatten" the C * H * W values into a single vector
        ↪per image

    def test_flatten():

```

```

x = torch.arange(12).view(2, 1, 3, 2)
print('Before flattening: ', x)
print('After flattening: ', flatten(x))

test_flatten()

```

```

Before flattening: tensor([[[[ 0,  1],
      [ 2,  3],
      [ 4,  5]]],

      [[[ 6,  7],
      [ 8,  9],
      [10, 11]]]])
After flattening: tensor([[ 0,  1,  2,  3,  4,  5],
      [ 6,  7,  8,  9, 10, 11]])

```

### 2.1.1 Check Accuracy Function

```

[45]: import torch.nn.functional as F # useful stateless functions
def check_accuracy(loader, model):
    if loader.dataset.train:
        print('Checking accuracy on validation set')
    else:
        print('Checking accuracy on test set')
    num_correct = 0
    num_samples = 0
    model.eval() # set model to evaluation mode
    with torch.no_grad():
        for x, y in loader:
            x = x.to(device=device, dtype=dtype) # move to device, e.g. GPU
            y = y.to(device=device, dtype=torch.long)
            scores = model(x)
            _, preds = scores.max(1)
            num_correct += (preds == y).sum()
            num_samples += preds.size(0)
        acc = float(num_correct) / num_samples
        print('Got %d / %d correct (%.2f)' % (num_correct, num_samples, 100 *
↪acc))
    return 100 * acc

```

### 2.1.2 Training Loop

```

[46]: def train(model, optimizer, epochs=1):
    """
    Train a model on CIFAR-10 using the PyTorch Module API.

```

*Inputs:*

- *model*: A PyTorch Module giving the model to train.
- *optimizer*: An Optimizer object we will use to train the model
- *epochs*: (Optional) A Python integer giving the number of epochs to train\_

↪ *for*

*Returns: Nothing, but prints model accuracies during training.*

```

"""
model = model.to(device=device) # move the model parameters to CPU/GPU
acc_max = 0
for e in range(epochs):
    for t, (x, y) in enumerate(loader_train):

        model.train() # put model to training mode
        x = x.to(device=device, dtype=dtype) # move to device, e.g. GPU
        y = y.to(device=device, dtype=torch.long)

        scores = model(x)
        loss = F.cross_entropy(scores, y)

        # Zero out all of the gradients for the variables which the_
↪ optimizer
        # will update.
        optimizer.zero_grad()

        # This is the backwards pass: compute the gradient of the loss with
        # respect to each parameter of the model.
        loss.backward()

        # Actually update the parameters of the model using the gradients
        # computed by the backwards pass.
        optimizer.step()

        if t % print_every == 0:
            print('Epoch %d, Iteration %d, loss = %.4f' % (e, t, loss.
↪ item()))

            acc = check_accuracy(loader_val, model)
            if acc >= acc_max:
                acc_max = acc
            print()
print("Maximum accuracy attained: ", acc_max)

```

```

[47]: # We need to wrap `flatten` function in a module in order to stack it
      # in nn.Sequential
      class Flatten(nn.Module):
          def forward(self, x):
              return flatten(x)

```

## 2.2 Vanilla CNN; No Attention

We implement the vanilla architecture for you here. Do not modify the architecture. You will use the same architecture in the following parts. Do not modify the hyper-parameters.

```
[94]: channel_1 = 64
channel_2 = 32
learning_rate = 1e-3
num_classes = 10

model = nn.Sequential(
    nn.Conv2d(3, channel_1, 3, padding=1, stride=1),
    nn.ReLU(),
    nn.Conv2d(channel_1, channel_2, 3, padding=1),
    nn.ReLU(),
    Flatten(),
    nn.Linear(channel_2*32*32, num_classes),
)

optimizer = optim.Adam(model.parameters(), lr=learning_rate)

train(model, optimizer, epochs=1)
```

```
Epoch 0, Iteration 0, loss = 2.2902
Checking accuracy on validation set
Got 119 / 1000 correct (11.90)
```

```
Epoch 0, Iteration 100, loss = 1.5537
Checking accuracy on validation set
Got 451 / 1000 correct (45.10)
```

```
Epoch 0, Iteration 200, loss = 1.5256
Checking accuracy on validation set
Got 473 / 1000 correct (47.30)
```

```
Epoch 0, Iteration 300, loss = 1.3774
Checking accuracy on validation set
Got 488 / 1000 correct (48.80)
```

```
Epoch 0, Iteration 400, loss = 1.0604
Checking accuracy on validation set
Got 535 / 1000 correct (53.50)
```

```
Epoch 0, Iteration 500, loss = 1.3256
Checking accuracy on validation set
Got 528 / 1000 correct (52.80)
```

```
Epoch 0, Iteration 600, loss = 1.2408
Checking accuracy on validation set
Got 570 / 1000 correct (57.00)
```

```
Epoch 0, Iteration 700, loss = 1.1947
Checking accuracy on validation set
Got 601 / 1000 correct (60.10)
```

```
Maximum accuracy attained: 60.099999999999994
```

### 2.3 Test set – run this only once

Now we test our model on the test set . Think about how this compares to your validation set accuracy. You should be able to see atleast 55% accuracy

```
[95]: vanillaModel = model
      check_accuracy(loader_test, vanillaModel)
```

```
Checking accuracy on test set
Got 5952 / 10000 correct (59.52)
```

```
[95]: 59.519999999999996
```

### 2.4 Part II Self-Attention

In the next section, you will implement an Attention layer which you will then use within a convnet architecture defined above for cifar 10 classification task.

A self-attention layer is formulated as following:

Input:  $X$  of shape  $(H \times W, C)$

Query, key, value linear transforms are  $W_Q, W_K, W_V$ , of shape  $(C, C)$ . We implement these linear transforms as 1x1 convolutional layers of the same dimensions.

$XW_Q, XW_K, XW_V$ , represent the output volumes when input  $X$  is passed through the transforms.

Self-Attention is given by the formula:  $Attention(X) = X + Softmax(\frac{XW_Q(XW_K)^T}{\sqrt{C}})XW_V$

#### 2.4.1 Inline Question 1: Self-Attention is equivalent to which of the following: (5 points)

1. K-means clustering
2. Non-local means
3. Residual Block
4. Gaussian Blurring

Your Answer: 2

## 2.4.2 Here you implement the Attention module, and run it in the next section (40 points)

```
[84]: # Initialize the attention module as a nn.Module subclass
class Attention(nn.Module):
    def __init__(self, in_channels):
        super().__init__()

        # TODO: Implement the Key, Query and Value linear transforms as 1x1
        ↪ convolutional layers
        # Hint: channel size remains constant throughout
        self.conv_query = nn.Conv2d(in_channels, in_channels, 1)
        self.conv_key = nn.Conv2d(in_channels, in_channels, 1)
        self.conv_value = nn.Conv2d(in_channels, in_channels, 1)

    def forward(self, x):
        N, C, H, W = x.shape

        # TODO: Pass the input through conv_query, reshape the output volume to
        ↪ (N, C, H*W)
        q = self.conv_query(x)
        q = torch.reshape(q, (N, C, -1))
        # TODO: Pass the input through conv_key, reshape the output volume to
        ↪ (N, C, H*W)
        k = self.conv_key(x)
        k = torch.reshape(k, (N, C, -1))
        # TODO: Pass the input through conv_value, reshape the output volume to
        ↪ (N, C, H*W)
        v = self.conv_value(x)
        v = torch.reshape(v, (N, C, -1))
        # TODO: Implement the above formula for attention using q, k, v, C
        # NOTE: The X in the formula is already added for you in the return line
        q = torch.transpose(q, 1, 2)
        att_distrib = torch.matmul(q, k)
        att_distrib = att_distrib / torch.sqrt(torch.tensor([C], device=device))
        att_distrib = torch.softmax(att_distrib, dim=1)
        # print(att_distrib.shape, torch.Tensor([C]))
        attention = torch.matmul(att_distrib, torch.transpose(v, 1, 2)) # (N,
        ↪ H*W, C)
        attention = torch.transpose(attention, 1, 2) # (N, C, H*W)
        # Reshape the output to (N, C, H, W) before adding to the input volume
        attention = attention.reshape(N, C, H, W)
        return x + attention
```

## 2.5 Single Attention Block: Early attention; After the first conv layer. (10 points)

```
[85]: channel_1 = 64
channel_2 = 32
learning_rate = 1e-3

# TODO: Use the above Attention module after the first Convolutional layer.
# Essentially the architecture should be
↳ [Conv->Relu->Attention->Relu->Conv->Relu->Linear]

model = nn.Sequential(
    nn.Conv2d(3, channel_1, 3, padding=1, stride=1),
    nn.ReLU(),
    Attention(channel_1),
    nn.ReLU(),
    nn.Conv2d(channel_1, channel_2, 3, padding=1, stride=1),
    nn.ReLU(),
    Flatten(),
    nn.Linear(channel_2*32*32, num_classes),
)

optimizer = optim.Adam(model.parameters(), lr=learning_rate)

train(model, optimizer, epochs=10)
```

```
Epoch 0, Iteration 0, loss = 2.2777
Checking accuracy on validation set
Got 145 / 1000 correct (14.50)
```

```
Epoch 0, Iteration 100, loss = 1.7835
Checking accuracy on validation set
Got 421 / 1000 correct (42.10)
```

```
Epoch 0, Iteration 200, loss = 1.3915
Checking accuracy on validation set
Got 451 / 1000 correct (45.10)
```

```
Epoch 0, Iteration 300, loss = 1.5871
Checking accuracy on validation set
Got 510 / 1000 correct (51.00)
```

```
Epoch 0, Iteration 400, loss = 1.4586
Checking accuracy on validation set
Got 529 / 1000 correct (52.90)
```

```
Epoch 0, Iteration 500, loss = 1.1199
```



Checking accuracy on validation set  
Got 506 / 1000 correct (50.60)

Epoch 0, Iteration 600, loss = 1.1241  
Checking accuracy on validation set  
Got 541 / 1000 correct (54.10)

Epoch 0, Iteration 700, loss = 1.0821  
Checking accuracy on validation set  
Got 552 / 1000 correct (55.20)

Epoch 1, Iteration 0, loss = 1.1126  
Checking accuracy on validation set  
Got 583 / 1000 correct (58.30)

Epoch 1, Iteration 100, loss = 1.2216  
Checking accuracy on validation set  
Got 580 / 1000 correct (58.00)

Epoch 1, Iteration 200, loss = 1.0497  
Checking accuracy on validation set  
Got 580 / 1000 correct (58.00)

Epoch 1, Iteration 300, loss = 1.1872  
Checking accuracy on validation set  
Got 615 / 1000 correct (61.50)

Epoch 1, Iteration 400, loss = 1.0578  
Checking accuracy on validation set  
Got 609 / 1000 correct (60.90)

Epoch 1, Iteration 500, loss = 0.9512  
Checking accuracy on validation set  
Got 626 / 1000 correct (62.60)

Epoch 1, Iteration 600, loss = 0.9448  
Checking accuracy on validation set  
Got 620 / 1000 correct (62.00)

Epoch 1, Iteration 700, loss = 1.0367  
Checking accuracy on validation set  
Got 620 / 1000 correct (62.00)

Epoch 2, Iteration 0, loss = 0.9432  
Checking accuracy on validation set  
Got 618 / 1000 correct (61.80)

Epoch 2, Iteration 100, loss = 1.2311

Checking accuracy on validation set  
Got 644 / 1000 correct (64.40)

Epoch 2, Iteration 200, loss = 1.0270  
Checking accuracy on validation set  
Got 625 / 1000 correct (62.50)

Epoch 2, Iteration 300, loss = 0.6286  
Checking accuracy on validation set  
Got 636 / 1000 correct (63.60)

Epoch 2, Iteration 400, loss = 0.7940  
Checking accuracy on validation set  
Got 629 / 1000 correct (62.90)

Epoch 2, Iteration 500, loss = 0.8733  
Checking accuracy on validation set  
Got 631 / 1000 correct (63.10)

Epoch 2, Iteration 600, loss = 0.7626  
Checking accuracy on validation set  
Got 656 / 1000 correct (65.60)

Epoch 2, Iteration 700, loss = 0.9407  
Checking accuracy on validation set  
Got 651 / 1000 correct (65.10)

Epoch 3, Iteration 0, loss = 0.6241  
Checking accuracy on validation set  
Got 646 / 1000 correct (64.60)

Epoch 3, Iteration 100, loss = 0.5352  
Checking accuracy on validation set  
Got 649 / 1000 correct (64.90)

Epoch 3, Iteration 200, loss = 0.7777  
Checking accuracy on validation set  
Got 640 / 1000 correct (64.00)

Epoch 3, Iteration 300, loss = 0.5430  
Checking accuracy on validation set  
Got 653 / 1000 correct (65.30)

Epoch 3, Iteration 400, loss = 0.7170  
Checking accuracy on validation set  
Got 648 / 1000 correct (64.80)

Epoch 3, Iteration 500, loss = 0.7274

Checking accuracy on validation set  
Got 655 / 1000 correct (65.50)

Epoch 3, Iteration 600, loss = 0.6654  
Checking accuracy on validation set  
Got 641 / 1000 correct (64.10)

Epoch 3, Iteration 700, loss = 0.7762  
Checking accuracy on validation set  
Got 656 / 1000 correct (65.60)

Epoch 4, Iteration 0, loss = 0.4312  
Checking accuracy on validation set  
Got 668 / 1000 correct (66.80)

Epoch 4, Iteration 100, loss = 0.5915  
Checking accuracy on validation set  
Got 650 / 1000 correct (65.00)

Epoch 4, Iteration 200, loss = 0.6140  
Checking accuracy on validation set  
Got 644 / 1000 correct (64.40)

Epoch 4, Iteration 300, loss = 0.6151  
Checking accuracy on validation set  
Got 640 / 1000 correct (64.00)

Epoch 4, Iteration 400, loss = 0.4157  
Checking accuracy on validation set  
Got 649 / 1000 correct (64.90)

Epoch 4, Iteration 500, loss = 0.5189  
Checking accuracy on validation set  
Got 648 / 1000 correct (64.80)

Epoch 4, Iteration 600, loss = 0.8579  
Checking accuracy on validation set  
Got 653 / 1000 correct (65.30)

Epoch 4, Iteration 700, loss = 0.6421  
Checking accuracy on validation set  
Got 654 / 1000 correct (65.40)

Epoch 5, Iteration 0, loss = 0.3781  
Checking accuracy on validation set  
Got 647 / 1000 correct (64.70)

Epoch 5, Iteration 100, loss = 0.4149

Checking accuracy on validation set  
Got 645 / 1000 correct (64.50)

Epoch 5, Iteration 200, loss = 0.5388  
Checking accuracy on validation set  
Got 637 / 1000 correct (63.70)

Epoch 5, Iteration 300, loss = 0.3386  
Checking accuracy on validation set  
Got 653 / 1000 correct (65.30)

Epoch 5, Iteration 400, loss = 0.2910  
Checking accuracy on validation set  
Got 664 / 1000 correct (66.40)

Epoch 5, Iteration 500, loss = 0.4093  
Checking accuracy on validation set  
Got 630 / 1000 correct (63.00)

Epoch 5, Iteration 600, loss = 0.3622  
Checking accuracy on validation set  
Got 646 / 1000 correct (64.60)

Epoch 5, Iteration 700, loss = 0.3400  
Checking accuracy on validation set  
Got 632 / 1000 correct (63.20)

Epoch 6, Iteration 0, loss = 0.1573  
Checking accuracy on validation set  
Got 650 / 1000 correct (65.00)

Epoch 6, Iteration 100, loss = 0.2719  
Checking accuracy on validation set  
Got 652 / 1000 correct (65.20)

Epoch 6, Iteration 200, loss = 0.2456  
Checking accuracy on validation set  
Got 638 / 1000 correct (63.80)

Epoch 6, Iteration 300, loss = 0.2910  
Checking accuracy on validation set  
Got 647 / 1000 correct (64.70)

Epoch 6, Iteration 400, loss = 0.2731  
Checking accuracy on validation set  
Got 644 / 1000 correct (64.40)

Epoch 6, Iteration 500, loss = 0.3571

Checking accuracy on validation set  
Got 623 / 1000 correct (62.30)

Epoch 6, Iteration 600, loss = 0.3025  
Checking accuracy on validation set  
Got 626 / 1000 correct (62.60)

Epoch 6, Iteration 700, loss = 0.2089  
Checking accuracy on validation set  
Got 622 / 1000 correct (62.20)

Epoch 7, Iteration 0, loss = 0.2481  
Checking accuracy on validation set  
Got 652 / 1000 correct (65.20)

Epoch 7, Iteration 100, loss = 0.1682  
Checking accuracy on validation set  
Got 647 / 1000 correct (64.70)

Epoch 7, Iteration 200, loss = 0.2243  
Checking accuracy on validation set  
Got 650 / 1000 correct (65.00)

Epoch 7, Iteration 300, loss = 0.1850  
Checking accuracy on validation set  
Got 636 / 1000 correct (63.60)

Epoch 7, Iteration 400, loss = 0.1420  
Checking accuracy on validation set  
Got 633 / 1000 correct (63.30)

Epoch 7, Iteration 500, loss = 0.1992  
Checking accuracy on validation set  
Got 638 / 1000 correct (63.80)

Epoch 7, Iteration 600, loss = 0.2456  
Checking accuracy on validation set  
Got 639 / 1000 correct (63.90)

Epoch 7, Iteration 700, loss = 0.2856  
Checking accuracy on validation set  
Got 629 / 1000 correct (62.90)

Epoch 8, Iteration 0, loss = 0.0965  
Checking accuracy on validation set  
Got 628 / 1000 correct (62.80)

Epoch 8, Iteration 100, loss = 0.1769

Checking accuracy on validation set  
Got 644 / 1000 correct (64.40)

Epoch 8, Iteration 200, loss = 0.0863  
Checking accuracy on validation set  
Got 638 / 1000 correct (63.80)

Epoch 8, Iteration 300, loss = 0.2237  
Checking accuracy on validation set  
Got 623 / 1000 correct (62.30)

Epoch 8, Iteration 400, loss = 0.3130  
Checking accuracy on validation set  
Got 617 / 1000 correct (61.70)

Epoch 8, Iteration 500, loss = 0.1709  
Checking accuracy on validation set  
Got 630 / 1000 correct (63.00)

Epoch 8, Iteration 600, loss = 0.0779  
Checking accuracy on validation set  
Got 614 / 1000 correct (61.40)

Epoch 8, Iteration 700, loss = 0.2981  
Checking accuracy on validation set  
Got 627 / 1000 correct (62.70)

Epoch 9, Iteration 0, loss = 0.0837  
Checking accuracy on validation set  
Got 626 / 1000 correct (62.60)

Epoch 9, Iteration 100, loss = 0.0700  
Checking accuracy on validation set  
Got 638 / 1000 correct (63.80)

Epoch 9, Iteration 200, loss = 0.1240  
Checking accuracy on validation set  
Got 634 / 1000 correct (63.40)

Epoch 9, Iteration 300, loss = 0.1387  
Checking accuracy on validation set  
Got 624 / 1000 correct (62.40)

Epoch 9, Iteration 400, loss = 0.1412  
Checking accuracy on validation set  
Got 631 / 1000 correct (63.10)

Epoch 9, Iteration 500, loss = 0.1931

```
Checking accuracy on validation set
Got 629 / 1000 correct (62.90)
```

```
Epoch 9, Iteration 600, loss = 0.1681
Checking accuracy on validation set
Got 630 / 1000 correct (63.00)
```

```
Epoch 9, Iteration 700, loss = 0.1472
Checking accuracy on validation set
Got 625 / 1000 correct (62.50)
```

```
Maximum accuracy attained: 66.8
```

## 2.6 Test set – run this only once

Now we test our model on the test set . Think about how this compares to your validation set accuracy. You should see improvement of about 2-3% over the vanilla convnet model. \* Use this part to tune your Attention module and then move on to the next parts. \*

```
[86]: earlyAttention = model
      check_accuracy(loader_test, earlyAttention)
```

```
Checking accuracy on test set
Got 6112 / 10000 correct (61.12)
```

```
[86]: 61.12
```

## 2.7 Single Attention Block: Late attention; After the second conv layer. (10 points)

```
[87]: channel_1 = 64
      channel_2 = 32
      learning_rate = 1e-3

      # TODO: Use the above Attention module after the Second Convolutional layer.
      # Essentially the architecture should be
      ↪ [Conv->Relu->Conv->Relu->Attention->Relu->Linear]

      model = nn.Sequential(
          nn.Conv2d(3, channel_1, 3, padding=1, stride=1),
          nn.ReLU(),
          nn.Conv2d(channel_1, channel_2, 3, padding=1, stride=1),
          nn.ReLU(),
          Attention(channel_2),
          nn.ReLU(),
          Flatten(),
          nn.Linear(channel_2*32*32, num_classes),
      )
```

```
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

train(model, optimizer, epochs=10)
```

Epoch 0, Iteration 0, loss = 2.3091  
Checking accuracy on validation set  
Got 125 / 1000 correct (12.50)

Epoch 0, Iteration 100, loss = 1.5839  
Checking accuracy on validation set  
Got 438 / 1000 correct (43.80)

Epoch 0, Iteration 200, loss = 1.9136  
Checking accuracy on validation set  
Got 497 / 1000 correct (49.70)

Epoch 0, Iteration 300, loss = 1.3352  
Checking accuracy on validation set  
Got 527 / 1000 correct (52.70)

Epoch 0, Iteration 400, loss = 1.3180  
Checking accuracy on validation set  
Got 532 / 1000 correct (53.20)

Epoch 0, Iteration 500, loss = 1.3627  
Checking accuracy on validation set  
Got 532 / 1000 correct (53.20)

Epoch 0, Iteration 600, loss = 1.4136  
Checking accuracy on validation set  
Got 554 / 1000 correct (55.40)

Epoch 0, Iteration 700, loss = 1.3577  
Checking accuracy on validation set  
Got 567 / 1000 correct (56.70)

Epoch 1, Iteration 0, loss = 1.3397  
Checking accuracy on validation set  
Got 557 / 1000 correct (55.70)

Epoch 1, Iteration 100, loss = 1.0281  
Checking accuracy on validation set  
Got 587 / 1000 correct (58.70)

Epoch 1, Iteration 200, loss = 0.8183  
Checking accuracy on validation set



Got 574 / 1000 correct (57.40)

Epoch 1, Iteration 300, loss = 0.9360  
Checking accuracy on validation set  
Got 586 / 1000 correct (58.60)

Epoch 1, Iteration 400, loss = 0.9947  
Checking accuracy on validation set  
Got 602 / 1000 correct (60.20)

Epoch 1, Iteration 500, loss = 0.8387  
Checking accuracy on validation set  
Got 607 / 1000 correct (60.70)

Epoch 1, Iteration 600, loss = 0.8237  
Checking accuracy on validation set  
Got 594 / 1000 correct (59.40)

Epoch 1, Iteration 700, loss = 0.9452  
Checking accuracy on validation set  
Got 585 / 1000 correct (58.50)

Epoch 2, Iteration 0, loss = 1.2590  
Checking accuracy on validation set  
Got 610 / 1000 correct (61.00)

Epoch 2, Iteration 100, loss = 0.9441  
Checking accuracy on validation set  
Got 589 / 1000 correct (58.90)

Epoch 2, Iteration 200, loss = 1.1428  
Checking accuracy on validation set  
Got 598 / 1000 correct (59.80)

Epoch 2, Iteration 300, loss = 0.7899  
Checking accuracy on validation set  
Got 585 / 1000 correct (58.50)

Epoch 2, Iteration 400, loss = 0.9110  
Checking accuracy on validation set  
Got 617 / 1000 correct (61.70)

Epoch 2, Iteration 500, loss = 0.9792  
Checking accuracy on validation set  
Got 605 / 1000 correct (60.50)

Epoch 2, Iteration 600, loss = 1.2413  
Checking accuracy on validation set

Got 620 / 1000 correct (62.00)

Epoch 2, Iteration 700, loss = 0.9184  
Checking accuracy on validation set  
Got 630 / 1000 correct (63.00)

Epoch 3, Iteration 0, loss = 0.7756  
Checking accuracy on validation set  
Got 619 / 1000 correct (61.90)

Epoch 3, Iteration 100, loss = 0.7184  
Checking accuracy on validation set  
Got 624 / 1000 correct (62.40)

Epoch 3, Iteration 200, loss = 0.6576  
Checking accuracy on validation set  
Got 612 / 1000 correct (61.20)

Epoch 3, Iteration 300, loss = 0.7145  
Checking accuracy on validation set  
Got 611 / 1000 correct (61.10)

Epoch 3, Iteration 400, loss = 0.9338  
Checking accuracy on validation set  
Got 632 / 1000 correct (63.20)

Epoch 3, Iteration 500, loss = 0.9933  
Checking accuracy on validation set  
Got 630 / 1000 correct (63.00)

Epoch 3, Iteration 600, loss = 0.5983  
Checking accuracy on validation set  
Got 618 / 1000 correct (61.80)

Epoch 3, Iteration 700, loss = 0.7197  
Checking accuracy on validation set  
Got 619 / 1000 correct (61.90)

Epoch 4, Iteration 0, loss = 0.6874  
Checking accuracy on validation set  
Got 619 / 1000 correct (61.90)

Epoch 4, Iteration 100, loss = 0.5170  
Checking accuracy on validation set  
Got 616 / 1000 correct (61.60)

Epoch 4, Iteration 200, loss = 0.4866  
Checking accuracy on validation set

Got 613 / 1000 correct (61.30)

Epoch 4, Iteration 300, loss = 1.1615  
Checking accuracy on validation set  
Got 611 / 1000 correct (61.10)

Epoch 4, Iteration 400, loss = 0.8170  
Checking accuracy on validation set  
Got 624 / 1000 correct (62.40)

Epoch 4, Iteration 500, loss = 1.0931  
Checking accuracy on validation set  
Got 595 / 1000 correct (59.50)

Epoch 4, Iteration 600, loss = 0.7278  
Checking accuracy on validation set  
Got 639 / 1000 correct (63.90)

Epoch 4, Iteration 700, loss = 0.5860  
Checking accuracy on validation set  
Got 622 / 1000 correct (62.20)

Epoch 5, Iteration 0, loss = 0.4848  
Checking accuracy on validation set  
Got 619 / 1000 correct (61.90)

Epoch 5, Iteration 100, loss = 0.4473  
Checking accuracy on validation set  
Got 610 / 1000 correct (61.00)

Epoch 5, Iteration 200, loss = 0.4268  
Checking accuracy on validation set  
Got 612 / 1000 correct (61.20)

Epoch 5, Iteration 300, loss = 0.6967  
Checking accuracy on validation set  
Got 616 / 1000 correct (61.60)

Epoch 5, Iteration 400, loss = 0.7446  
Checking accuracy on validation set  
Got 608 / 1000 correct (60.80)

Epoch 5, Iteration 500, loss = 0.5288  
Checking accuracy on validation set  
Got 618 / 1000 correct (61.80)

Epoch 5, Iteration 600, loss = 0.8227  
Checking accuracy on validation set

Got 612 / 1000 correct (61.20)

Epoch 5, Iteration 700, loss = 0.4823  
Checking accuracy on validation set  
Got 625 / 1000 correct (62.50)

Epoch 6, Iteration 0, loss = 0.5073  
Checking accuracy on validation set  
Got 629 / 1000 correct (62.90)

Epoch 6, Iteration 100, loss = 0.4172  
Checking accuracy on validation set  
Got 612 / 1000 correct (61.20)

Epoch 6, Iteration 200, loss = 0.2946  
Checking accuracy on validation set  
Got 617 / 1000 correct (61.70)

Epoch 6, Iteration 300, loss = 0.3361  
Checking accuracy on validation set  
Got 607 / 1000 correct (60.70)

Epoch 6, Iteration 400, loss = 0.5107  
Checking accuracy on validation set  
Got 590 / 1000 correct (59.00)

Epoch 6, Iteration 500, loss = 0.4370  
Checking accuracy on validation set  
Got 618 / 1000 correct (61.80)

Epoch 6, Iteration 600, loss = 0.6390  
Checking accuracy on validation set  
Got 597 / 1000 correct (59.70)

Epoch 6, Iteration 700, loss = 0.4263  
Checking accuracy on validation set  
Got 599 / 1000 correct (59.90)

Epoch 7, Iteration 0, loss = 0.3595  
Checking accuracy on validation set  
Got 595 / 1000 correct (59.50)

Epoch 7, Iteration 100, loss = 0.2347  
Checking accuracy on validation set  
Got 600 / 1000 correct (60.00)

Epoch 7, Iteration 200, loss = 0.4223  
Checking accuracy on validation set

Got 598 / 1000 correct (59.80)

Epoch 7, Iteration 300, loss = 0.2486  
Checking accuracy on validation set  
Got 592 / 1000 correct (59.20)

Epoch 7, Iteration 400, loss = 0.5701  
Checking accuracy on validation set  
Got 587 / 1000 correct (58.70)

Epoch 7, Iteration 500, loss = 0.3920  
Checking accuracy on validation set  
Got 591 / 1000 correct (59.10)

Epoch 7, Iteration 600, loss = 0.6804  
Checking accuracy on validation set  
Got 607 / 1000 correct (60.70)

Epoch 7, Iteration 700, loss = 0.4118  
Checking accuracy on validation set  
Got 595 / 1000 correct (59.50)

Epoch 8, Iteration 0, loss = 0.3907  
Checking accuracy on validation set  
Got 614 / 1000 correct (61.40)

Epoch 8, Iteration 100, loss = 0.2254  
Checking accuracy on validation set  
Got 596 / 1000 correct (59.60)

Epoch 8, Iteration 200, loss = 0.2503  
Checking accuracy on validation set  
Got 598 / 1000 correct (59.80)

Epoch 8, Iteration 300, loss = 0.3949  
Checking accuracy on validation set  
Got 578 / 1000 correct (57.80)

Epoch 8, Iteration 400, loss = 0.3272  
Checking accuracy on validation set  
Got 581 / 1000 correct (58.10)

Epoch 8, Iteration 500, loss = 0.4342  
Checking accuracy on validation set  
Got 594 / 1000 correct (59.40)

Epoch 8, Iteration 600, loss = 0.3267  
Checking accuracy on validation set

```
Got 591 / 1000 correct (59.10)

Epoch 8, Iteration 700, loss = 0.4204
Checking accuracy on validation set
Got 602 / 1000 correct (60.20)

Epoch 9, Iteration 0, loss = 0.2279
Checking accuracy on validation set
Got 575 / 1000 correct (57.50)

Epoch 9, Iteration 100, loss = 0.2174
Checking accuracy on validation set
Got 583 / 1000 correct (58.30)

Epoch 9, Iteration 200, loss = 0.1031
Checking accuracy on validation set
Got 584 / 1000 correct (58.40)

Epoch 9, Iteration 300, loss = 0.2592
Checking accuracy on validation set
Got 583 / 1000 correct (58.30)

Epoch 9, Iteration 400, loss = 0.2464
Checking accuracy on validation set
Got 578 / 1000 correct (57.80)

Epoch 9, Iteration 500, loss = 0.2619
Checking accuracy on validation set
Got 567 / 1000 correct (56.70)

Epoch 9, Iteration 600, loss = 0.3589
Checking accuracy on validation set
Got 572 / 1000 correct (57.20)

Epoch 9, Iteration 700, loss = 0.1551
Checking accuracy on validation set
Got 577 / 1000 correct (57.70)

Maximum accuracy attained: 63.9
```

## 2.8 Test set – run this only once

Now we test our model on the test set . Think about how this compares to your validation set accuracy.

```
[88]: lateAttention = model
      check_accuracy(loader_test, lateAttention)
```

Checking accuracy on test set  
Got 5831 / 10000 correct (58.31)

[88]: 58.309999999999995

### 2.8.1 Inline Question 2: Provide one example each of usage of self-attention and attention in computer vision. Explain the difference between the two. (5 points)

Your Answer: \* Usage of self-attention: The Vision Transformer (ViT) leverages self-attention to process images as sequences of patches, enabling it to capture long-range dependencies and interactions between different parts of the image. This approach provides an alternative to traditional CNNs in object classification, offering competitive performance on various computer vision tasks.

- Usage of attention: The Squeeze-and-Excitation Network (SENet) enhances CNNs by recalibrating channel-wise feature responses using attention method. This allows the network to focus on the most relevant parts of the input sequence, improving performance and global feature representation.
- Differences:
  - Self-attention is typically part of a layer (like in ViT) capturing non-local dependencies by computing relationships between all pairs of elements in an input. It can be used very frequently in a network. It models within-sequence relationships mainly. The weights are computed from the input sequence directly, and hence different inputs correspond to different weights. It improves performance by capturing dependencies within the input sequence more effectively by doing dot-products using queries and keys for the weights.
  - Attention is typically an individual module that concatenates encoder and decoder, serving the function of capturing the sequence-to-sequence relationships in the network. It's often used only a few times to concatenate two modules (like an encoder and a decoder) in a network. And hence, it models sequence-to-sequence relationships mainly. The weights in the MLP are learnt, and hence different inputs correspond to the same group of learnt weights. Attention improves performance by emphasizing important features based on the learnt weights indicating relative importance.

### 2.9 Double Attention Blocks: After conv layers 1 and 2 (10 points)

```
[89]: channel_1 = 64
channel_2 = 32
learning_rate = 1e-3

# TODO: Use the above Attention module after the Second Convolutional layer.
# Essentially the architecture should be
→ [Conv->Relu->Attention->Relu->Conv->Relu->Attention->Relu->Linear]

model = nn.Sequential(
    nn.Conv2d(3, channel_1, 3, padding=1, stride=1),
    nn.ReLU(),
    Attention(channel_1),
```

```

    nn.ReLU(),
    nn.Conv2d(channel_1, channel_2, 3, padding=1, stride=1),
    nn.ReLU(),
    Attention(channel_2),
    nn.ReLU(),
    Flatten(),
    nn.Linear(channel_2*32*32, num_classes),
)

optimizer = optim.Adam(model.parameters(), lr=learning_rate)

train(model, optimizer, epochs=10)

```

Epoch 0, Iteration 0, loss = 2.2909  
 Checking accuracy on validation set  
 Got 109 / 1000 correct (10.90)

Epoch 0, Iteration 100, loss = 1.9164  
 Checking accuracy on validation set  
 Got 322 / 1000 correct (32.20)

Epoch 0, Iteration 200, loss = 1.5259  
 Checking accuracy on validation set  
 Got 450 / 1000 correct (45.00)

Epoch 0, Iteration 300, loss = 1.3118  
 Checking accuracy on validation set  
 Got 467 / 1000 correct (46.70)

Epoch 0, Iteration 400, loss = 1.4503  
 Checking accuracy on validation set  
 Got 506 / 1000 correct (50.60)

Epoch 0, Iteration 500, loss = 1.3227  
 Checking accuracy on validation set  
 Got 515 / 1000 correct (51.50)

Epoch 0, Iteration 600, loss = 1.2445  
 Checking accuracy on validation set  
 Got 524 / 1000 correct (52.40)

Epoch 0, Iteration 700, loss = 1.3021  
 Checking accuracy on validation set  
 Got 532 / 1000 correct (53.20)

Epoch 1, Iteration 0, loss = 1.2930  
 Checking accuracy on validation set



Got 530 / 1000 correct (53.00)

Epoch 1, Iteration 100, loss = 1.3517  
Checking accuracy on validation set  
Got 542 / 1000 correct (54.20)

Epoch 1, Iteration 200, loss = 1.2945  
Checking accuracy on validation set  
Got 567 / 1000 correct (56.70)

Epoch 1, Iteration 300, loss = 1.3331  
Checking accuracy on validation set  
Got 538 / 1000 correct (53.80)

Epoch 1, Iteration 400, loss = 1.0128  
Checking accuracy on validation set  
Got 551 / 1000 correct (55.10)

Epoch 1, Iteration 500, loss = 1.3041  
Checking accuracy on validation set  
Got 530 / 1000 correct (53.00)

Epoch 1, Iteration 600, loss = 1.2753  
Checking accuracy on validation set  
Got 551 / 1000 correct (55.10)

Epoch 1, Iteration 700, loss = 1.1456  
Checking accuracy on validation set  
Got 576 / 1000 correct (57.60)

Epoch 2, Iteration 0, loss = 1.1693  
Checking accuracy on validation set  
Got 576 / 1000 correct (57.60)

Epoch 2, Iteration 100, loss = 1.0134  
Checking accuracy on validation set  
Got 587 / 1000 correct (58.70)

Epoch 2, Iteration 200, loss = 1.1225  
Checking accuracy on validation set  
Got 584 / 1000 correct (58.40)

Epoch 2, Iteration 300, loss = 1.3498  
Checking accuracy on validation set  
Got 588 / 1000 correct (58.80)

Epoch 2, Iteration 400, loss = 1.2768  
Checking accuracy on validation set

Got 596 / 1000 correct (59.60)

Epoch 2, Iteration 500, loss = 1.1820  
Checking accuracy on validation set  
Got 604 / 1000 correct (60.40)

Epoch 2, Iteration 600, loss = 1.4075  
Checking accuracy on validation set  
Got 617 / 1000 correct (61.70)

Epoch 2, Iteration 700, loss = 1.0997  
Checking accuracy on validation set  
Got 598 / 1000 correct (59.80)

Epoch 3, Iteration 0, loss = 0.9658  
Checking accuracy on validation set  
Got 604 / 1000 correct (60.40)

Epoch 3, Iteration 100, loss = 0.7417  
Checking accuracy on validation set  
Got 604 / 1000 correct (60.40)

Epoch 3, Iteration 200, loss = 0.9589  
Checking accuracy on validation set  
Got 589 / 1000 correct (58.90)

Epoch 3, Iteration 300, loss = 1.1804  
Checking accuracy on validation set  
Got 595 / 1000 correct (59.50)

Epoch 3, Iteration 400, loss = 1.1039  
Checking accuracy on validation set  
Got 622 / 1000 correct (62.20)

Epoch 3, Iteration 500, loss = 1.0229  
Checking accuracy on validation set  
Got 593 / 1000 correct (59.30)

Epoch 3, Iteration 600, loss = 0.9299  
Checking accuracy on validation set  
Got 612 / 1000 correct (61.20)

Epoch 3, Iteration 700, loss = 1.3686  
Checking accuracy on validation set  
Got 614 / 1000 correct (61.40)

Epoch 4, Iteration 0, loss = 0.6392  
Checking accuracy on validation set

Got 607 / 1000 correct (60.70)

Epoch 4, Iteration 100, loss = 0.8663  
Checking accuracy on validation set  
Got 588 / 1000 correct (58.80)

Epoch 4, Iteration 200, loss = 0.8482  
Checking accuracy on validation set  
Got 597 / 1000 correct (59.70)

Epoch 4, Iteration 300, loss = 0.8884  
Checking accuracy on validation set  
Got 590 / 1000 correct (59.00)

Epoch 4, Iteration 400, loss = 0.9475  
Checking accuracy on validation set  
Got 612 / 1000 correct (61.20)

Epoch 4, Iteration 500, loss = 1.0487  
Checking accuracy on validation set  
Got 615 / 1000 correct (61.50)

Epoch 4, Iteration 600, loss = 0.7489  
Checking accuracy on validation set  
Got 617 / 1000 correct (61.70)

Epoch 4, Iteration 700, loss = 0.8002  
Checking accuracy on validation set  
Got 596 / 1000 correct (59.60)

Epoch 5, Iteration 0, loss = 0.5266  
Checking accuracy on validation set  
Got 615 / 1000 correct (61.50)

Epoch 5, Iteration 100, loss = 0.6291  
Checking accuracy on validation set  
Got 620 / 1000 correct (62.00)

Epoch 5, Iteration 200, loss = 0.6703  
Checking accuracy on validation set  
Got 615 / 1000 correct (61.50)

Epoch 5, Iteration 300, loss = 0.5306  
Checking accuracy on validation set  
Got 592 / 1000 correct (59.20)

Epoch 5, Iteration 400, loss = 0.9859  
Checking accuracy on validation set

Got 611 / 1000 correct (61.10)

Epoch 5, Iteration 500, loss = 0.8990  
Checking accuracy on validation set  
Got 626 / 1000 correct (62.60)

Epoch 5, Iteration 600, loss = 0.7123  
Checking accuracy on validation set  
Got 637 / 1000 correct (63.70)

Epoch 5, Iteration 700, loss = 0.8093  
Checking accuracy on validation set  
Got 613 / 1000 correct (61.30)

Epoch 6, Iteration 0, loss = 0.4293  
Checking accuracy on validation set  
Got 633 / 1000 correct (63.30)

Epoch 6, Iteration 100, loss = 0.6259  
Checking accuracy on validation set  
Got 606 / 1000 correct (60.60)

Epoch 6, Iteration 200, loss = 0.5768  
Checking accuracy on validation set  
Got 621 / 1000 correct (62.10)

Epoch 6, Iteration 300, loss = 0.4714  
Checking accuracy on validation set  
Got 612 / 1000 correct (61.20)

Epoch 6, Iteration 400, loss = 0.4936  
Checking accuracy on validation set  
Got 615 / 1000 correct (61.50)

Epoch 6, Iteration 500, loss = 0.8145  
Checking accuracy on validation set  
Got 616 / 1000 correct (61.60)

Epoch 6, Iteration 600, loss = 0.5466  
Checking accuracy on validation set  
Got 601 / 1000 correct (60.10)

Epoch 6, Iteration 700, loss = 0.6788  
Checking accuracy on validation set  
Got 615 / 1000 correct (61.50)

Epoch 7, Iteration 0, loss = 0.5111  
Checking accuracy on validation set

Got 618 / 1000 correct (61.80)

Epoch 7, Iteration 100, loss = 0.3314  
Checking accuracy on validation set  
Got 602 / 1000 correct (60.20)

Epoch 7, Iteration 200, loss = 0.4677  
Checking accuracy on validation set  
Got 594 / 1000 correct (59.40)

Epoch 7, Iteration 300, loss = 0.4808  
Checking accuracy on validation set  
Got 608 / 1000 correct (60.80)

Epoch 7, Iteration 400, loss = 0.4332  
Checking accuracy on validation set  
Got 607 / 1000 correct (60.70)

Epoch 7, Iteration 500, loss = 0.5851  
Checking accuracy on validation set  
Got 599 / 1000 correct (59.90)

Epoch 7, Iteration 600, loss = 0.4782  
Checking accuracy on validation set  
Got 617 / 1000 correct (61.70)

Epoch 7, Iteration 700, loss = 0.4484  
Checking accuracy on validation set  
Got 614 / 1000 correct (61.40)

Epoch 8, Iteration 0, loss = 0.2564  
Checking accuracy on validation set  
Got 606 / 1000 correct (60.60)

Epoch 8, Iteration 100, loss = 0.2846  
Checking accuracy on validation set  
Got 610 / 1000 correct (61.00)

Epoch 8, Iteration 200, loss = 0.3426  
Checking accuracy on validation set  
Got 599 / 1000 correct (59.90)

Epoch 8, Iteration 300, loss = 0.3699  
Checking accuracy on validation set  
Got 589 / 1000 correct (58.90)

Epoch 8, Iteration 400, loss = 0.2428  
Checking accuracy on validation set

Got 605 / 1000 correct (60.50)

Epoch 8, Iteration 500, loss = 0.3250  
Checking accuracy on validation set  
Got 595 / 1000 correct (59.50)

Epoch 8, Iteration 600, loss = 0.4352  
Checking accuracy on validation set  
Got 587 / 1000 correct (58.70)

Epoch 8, Iteration 700, loss = 0.5713  
Checking accuracy on validation set  
Got 570 / 1000 correct (57.00)

Epoch 9, Iteration 0, loss = 0.2609  
Checking accuracy on validation set  
Got 594 / 1000 correct (59.40)

Epoch 9, Iteration 100, loss = 0.1942  
Checking accuracy on validation set  
Got 594 / 1000 correct (59.40)

Epoch 9, Iteration 200, loss = 0.2212  
Checking accuracy on validation set  
Got 589 / 1000 correct (58.90)

Epoch 9, Iteration 300, loss = 0.2576  
Checking accuracy on validation set  
Got 584 / 1000 correct (58.40)

Epoch 9, Iteration 400, loss = 0.2013  
Checking accuracy on validation set  
Got 601 / 1000 correct (60.10)

Epoch 9, Iteration 500, loss = 0.4081  
Checking accuracy on validation set  
Got 586 / 1000 correct (58.60)

Epoch 9, Iteration 600, loss = 0.4593  
Checking accuracy on validation set  
Got 584 / 1000 correct (58.40)

Epoch 9, Iteration 700, loss = 0.3191  
Checking accuracy on validation set  
Got 593 / 1000 correct (59.30)

Maximum accuracy attained: 63.7

## 2.10 Test set – run this only once

Now we test our model on the test set . Think about how this compares to your validation set accuracy.

```
[90]: vanillaModel = model
      check_accuracy(loader_test, vanillaModel)
```

Checking accuracy on test set  
Got 5862 / 10000 correct (58.62)

```
[90]: 58.620000000000005
```

## 2.11 Resnet with Attention

Now we will experiment with applying attention within the Resnet10 architecture that we implemented in Homework 2. Please note that for a deeper model such as Resnet we do not expect significant improvements in performance with Attention

## 2.12 Vanilla Resnet, No Attention

The architecture for Resnet is given below, please train it and evaluate it on the test set.

```
[57]: import torch
      import torch.nn as nn

      class ResNet(nn.Module):

          def __init__(self, block, layers, img_channels=3, num_classes=100,
      ↪batchnorm=False):
              super(ResNet, self).__init__() #layers = [1, 1, 1, 1]
              self.in_channels = 64
              self.conv1 = nn.Conv2d(img_channels, 64, kernel_size=7, stride=2,
      ↪padding=3)
              self.bn1 = nn.BatchNorm2d(64)
              self.relu = nn.ReLU()
              self.maxpool = nn.MaxPool2d(kernel_size=3, stride=2, padding=1)
              self.batchnorm = batchnorm
              self.layer1 = self.make_layer(block, layers[0], out_channels=64,
      ↪stride=1, batchnorm=batchnorm)
              self.layer2 = self.make_layer(block, layers[1], out_channels=128,
      ↪stride=1, batchnorm=batchnorm)
              self.layer3 = self.make_layer(block, layers[2], out_channels=256,
      ↪stride=1, batchnorm=batchnorm)
              self.layer4 = self.make_layer(block, layers[3], out_channels=512,
      ↪stride=2, batchnorm=batchnorm)

              self.averagepool = nn.AdaptiveAvgPool2d((1, 1))
```

```

        self.fc = nn.Linear(512, num_classes)

    def forward(self, x):

        x = self.conv1(x)
        if self.batchnorm:
            x = self.bn1(x)
        x = self.relu(x)
        x = self.maxpool(x)
        x = self.layer1(x)
        x = self.layer2(x)
        x = self.layer3(x)
        x = self.layer4(x)
        x = self.averagepool(x)
        x = x.reshape(x.shape[0], -1)
        x = x.reshape(x.shape[0], -1)
        x = self.fc(x)

    return x

    def make_layer(self, block, num_blocks, out_channels, stride,
        ↪batchnorm=False):
        downsampler = None
        layers = []
        if stride != 1 or self.in_channels != out_channels:
            downsampler = nn.Sequential(nn.Conv2d(self.in_channels,
        ↪out_channels, kernel_size = 1, stride = stride), nn.
        ↪BatchNorm2d(out_channels))

        layers.append(block(self.in_channels, out_channels, downsampler,
        ↪stride, batchnorm=batchnorm))

        self.in_channels = out_channels

        for i in range(num_blocks - 1):
            layers.append(block(self.in_channels, out_channels))

        return nn.Sequential(*layers)

class block(nn.Module):

```



```

    def __init__(self, in_channels, out_channels, downsampler = None, stride = 1, batchnorm=False):

        super(block, self).__init__()
        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size = 3, padding = 2)
        self.bn1 = nn.BatchNorm2d(out_channels)
        self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size = 3, stride = stride)
        self.bn2 = nn.BatchNorm2d(out_channels)
        self.downsampler = downsampler
        self.relu = nn.ReLU()
        self.batchnorm = batchnorm

    def forward(self, x):

        residual = x
        x = self.conv1(x)
        if self.batchnorm:
            x = self.bn1(x)
        x = self.relu(x)
        x = self.conv2(x)
        if self.batchnorm:
            x = self.bn2(x)
        x = self.relu(x)

        if self.downsampler:
            residual = self.downsampler(residual)

        return self.relu(residual + x)

def ResNet10(num_classes = 100, batchnorm= False):

    return ResNet(block, [1, 1, 1, 1], num_classes=num_classes, batchnorm=batchnorm)

```

## 2.13 Test set – run this only once

Now we test our model on the test set . Think about how this compares to your validation set accuracy.

```
[58]: learning_rate = 1e-3
```

```
model = ResNet10()
```

```
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

train(model, optimizer, epochs=10)

vanillaResnet = model
check_accuracy(loader_test, vanillaResnet)
```

Epoch 0, Iteration 0, loss = 4.7750  
Checking accuracy on validation set  
Got 105 / 1000 correct (10.50)

Epoch 0, Iteration 100, loss = 1.5069  
Checking accuracy on validation set  
Got 395 / 1000 correct (39.50)

Epoch 0, Iteration 200, loss = 1.3592  
Checking accuracy on validation set  
Got 449 / 1000 correct (44.90)

Epoch 0, Iteration 300, loss = 1.4675  
Checking accuracy on validation set  
Got 432 / 1000 correct (43.20)

Epoch 0, Iteration 400, loss = 1.2042  
Checking accuracy on validation set  
Got 471 / 1000 correct (47.10)

Epoch 0, Iteration 500, loss = 1.1853  
Checking accuracy on validation set  
Got 528 / 1000 correct (52.80)

Epoch 0, Iteration 600, loss = 1.1258  
Checking accuracy on validation set  
Got 571 / 1000 correct (57.10)

Epoch 0, Iteration 700, loss = 1.0300  
Checking accuracy on validation set  
Got 573 / 1000 correct (57.30)

Epoch 1, Iteration 0, loss = 1.0037  
Checking accuracy on validation set  
Got 577 / 1000 correct (57.70)

Epoch 1, Iteration 100, loss = 0.8076  
Checking accuracy on validation set  
Got 528 / 1000 correct (52.80)

Epoch 1, Iteration 200, loss = 0.9683  
Checking accuracy on validation set  
Got 623 / 1000 correct (62.30)

Epoch 1, Iteration 300, loss = 1.2502  
Checking accuracy on validation set  
Got 538 / 1000 correct (53.80)

Epoch 1, Iteration 400, loss = 1.0313  
Checking accuracy on validation set  
Got 621 / 1000 correct (62.10)

Epoch 1, Iteration 500, loss = 1.0359  
Checking accuracy on validation set  
Got 612 / 1000 correct (61.20)

Epoch 1, Iteration 600, loss = 1.1694  
Checking accuracy on validation set  
Got 628 / 1000 correct (62.80)

Epoch 1, Iteration 700, loss = 1.0434  
Checking accuracy on validation set  
Got 621 / 1000 correct (62.10)

Epoch 2, Iteration 0, loss = 0.9030  
Checking accuracy on validation set  
Got 631 / 1000 correct (63.10)

Epoch 2, Iteration 100, loss = 1.0406  
Checking accuracy on validation set  
Got 609 / 1000 correct (60.90)

Epoch 2, Iteration 200, loss = 0.7680  
Checking accuracy on validation set  
Got 640 / 1000 correct (64.00)

Epoch 2, Iteration 300, loss = 0.8748  
Checking accuracy on validation set  
Got 665 / 1000 correct (66.50)

Epoch 2, Iteration 400, loss = 0.9845  
Checking accuracy on validation set  
Got 658 / 1000 correct (65.80)

Epoch 2, Iteration 500, loss = 0.9294  
Checking accuracy on validation set  
Got 634 / 1000 correct (63.40)

Epoch 2, Iteration 600, loss = 1.1083  
Checking accuracy on validation set  
Got 662 / 1000 correct (66.20)

Epoch 2, Iteration 700, loss = 0.6598  
Checking accuracy on validation set  
Got 697 / 1000 correct (69.70)

Epoch 3, Iteration 0, loss = 0.8251  
Checking accuracy on validation set  
Got 674 / 1000 correct (67.40)

Epoch 3, Iteration 100, loss = 0.5967  
Checking accuracy on validation set  
Got 692 / 1000 correct (69.20)

Epoch 3, Iteration 200, loss = 0.7270  
Checking accuracy on validation set  
Got 680 / 1000 correct (68.00)

Epoch 3, Iteration 300, loss = 1.1079  
Checking accuracy on validation set  
Got 676 / 1000 correct (67.60)

Epoch 3, Iteration 400, loss = 0.6674  
Checking accuracy on validation set  
Got 704 / 1000 correct (70.40)

Epoch 3, Iteration 500, loss = 0.6718  
Checking accuracy on validation set  
Got 706 / 1000 correct (70.60)

Epoch 3, Iteration 600, loss = 0.8633  
Checking accuracy on validation set  
Got 707 / 1000 correct (70.70)

Epoch 3, Iteration 700, loss = 0.8723  
Checking accuracy on validation set  
Got 718 / 1000 correct (71.80)

Epoch 4, Iteration 0, loss = 0.6791  
Checking accuracy on validation set  
Got 683 / 1000 correct (68.30)

Epoch 4, Iteration 100, loss = 0.6927  
Checking accuracy on validation set  
Got 709 / 1000 correct (70.90)

Epoch 4, Iteration 200, loss = 0.5078  
Checking accuracy on validation set  
Got 706 / 1000 correct (70.60)

Epoch 4, Iteration 300, loss = 0.6036  
Checking accuracy on validation set  
Got 736 / 1000 correct (73.60)

Epoch 4, Iteration 400, loss = 0.7852  
Checking accuracy on validation set  
Got 712 / 1000 correct (71.20)

Epoch 4, Iteration 500, loss = 0.8443  
Checking accuracy on validation set  
Got 713 / 1000 correct (71.30)

Epoch 4, Iteration 600, loss = 0.6034  
Checking accuracy on validation set  
Got 730 / 1000 correct (73.00)

Epoch 4, Iteration 700, loss = 0.7621  
Checking accuracy on validation set  
Got 730 / 1000 correct (73.00)

Epoch 5, Iteration 0, loss = 0.5981  
Checking accuracy on validation set  
Got 726 / 1000 correct (72.60)

Epoch 5, Iteration 100, loss = 0.9903  
Checking accuracy on validation set  
Got 730 / 1000 correct (73.00)

Epoch 5, Iteration 200, loss = 0.6806  
Checking accuracy on validation set  
Got 739 / 1000 correct (73.90)

Epoch 5, Iteration 300, loss = 0.4643  
Checking accuracy on validation set  
Got 703 / 1000 correct (70.30)

Epoch 5, Iteration 400, loss = 0.6257  
Checking accuracy on validation set  
Got 756 / 1000 correct (75.60)

Epoch 5, Iteration 500, loss = 0.6085  
Checking accuracy on validation set  
Got 716 / 1000 correct (71.60)

Epoch 5, Iteration 600, loss = 0.4997  
Checking accuracy on validation set  
Got 727 / 1000 correct (72.70)

Epoch 5, Iteration 700, loss = 0.6201  
Checking accuracy on validation set  
Got 742 / 1000 correct (74.20)

Epoch 6, Iteration 0, loss = 0.5664  
Checking accuracy on validation set  
Got 735 / 1000 correct (73.50)

Epoch 6, Iteration 100, loss = 0.5474  
Checking accuracy on validation set  
Got 730 / 1000 correct (73.00)

Epoch 6, Iteration 200, loss = 0.6548  
Checking accuracy on validation set  
Got 735 / 1000 correct (73.50)

Epoch 6, Iteration 300, loss = 0.4773  
Checking accuracy on validation set  
Got 739 / 1000 correct (73.90)

Epoch 6, Iteration 400, loss = 0.4326  
Checking accuracy on validation set  
Got 736 / 1000 correct (73.60)

Epoch 6, Iteration 500, loss = 0.5066  
Checking accuracy on validation set  
Got 761 / 1000 correct (76.10)

Epoch 6, Iteration 600, loss = 0.4803  
Checking accuracy on validation set  
Got 750 / 1000 correct (75.00)

Epoch 6, Iteration 700, loss = 0.5982  
Checking accuracy on validation set  
Got 747 / 1000 correct (74.70)

Epoch 7, Iteration 0, loss = 0.5280  
Checking accuracy on validation set  
Got 742 / 1000 correct (74.20)

Epoch 7, Iteration 100, loss = 0.5425  
Checking accuracy on validation set  
Got 729 / 1000 correct (72.90)

Epoch 7, Iteration 200, loss = 0.2875  
Checking accuracy on validation set  
Got 764 / 1000 correct (76.40)

Epoch 7, Iteration 300, loss = 0.4327  
Checking accuracy on validation set  
Got 743 / 1000 correct (74.30)

Epoch 7, Iteration 400, loss = 0.4404  
Checking accuracy on validation set  
Got 734 / 1000 correct (73.40)

Epoch 7, Iteration 500, loss = 0.6656  
Checking accuracy on validation set  
Got 752 / 1000 correct (75.20)

Epoch 7, Iteration 600, loss = 0.7700  
Checking accuracy on validation set  
Got 762 / 1000 correct (76.20)

Epoch 7, Iteration 700, loss = 0.3530  
Checking accuracy on validation set  
Got 753 / 1000 correct (75.30)

Epoch 8, Iteration 0, loss = 0.4891  
Checking accuracy on validation set  
Got 745 / 1000 correct (74.50)

Epoch 8, Iteration 100, loss = 0.3552  
Checking accuracy on validation set  
Got 743 / 1000 correct (74.30)

Epoch 8, Iteration 200, loss = 0.5656  
Checking accuracy on validation set  
Got 738 / 1000 correct (73.80)

Epoch 8, Iteration 300, loss = 0.6213  
Checking accuracy on validation set  
Got 756 / 1000 correct (75.60)

Epoch 8, Iteration 400, loss = 0.6020  
Checking accuracy on validation set  
Got 733 / 1000 correct (73.30)

Epoch 8, Iteration 500, loss = 0.2505  
Checking accuracy on validation set  
Got 753 / 1000 correct (75.30)

Epoch 8, Iteration 600, loss = 0.5816  
Checking accuracy on validation set  
Got 770 / 1000 correct (77.00)

Epoch 8, Iteration 700, loss = 0.2314  
Checking accuracy on validation set  
Got 752 / 1000 correct (75.20)

Epoch 9, Iteration 0, loss = 0.4711  
Checking accuracy on validation set  
Got 761 / 1000 correct (76.10)

Epoch 9, Iteration 100, loss = 0.3250  
Checking accuracy on validation set  
Got 754 / 1000 correct (75.40)

Epoch 9, Iteration 200, loss = 0.3695  
Checking accuracy on validation set  
Got 762 / 1000 correct (76.20)

Epoch 9, Iteration 300, loss = 0.4426  
Checking accuracy on validation set  
Got 762 / 1000 correct (76.20)

Epoch 9, Iteration 400, loss = 0.3072  
Checking accuracy on validation set  
Got 764 / 1000 correct (76.40)

Epoch 9, Iteration 500, loss = 0.4218  
Checking accuracy on validation set  
Got 770 / 1000 correct (77.00)

Epoch 9, Iteration 600, loss = 0.4738  
Checking accuracy on validation set  
Got 769 / 1000 correct (76.90)

Epoch 9, Iteration 700, loss = 0.4158  
Checking accuracy on validation set  
Got 743 / 1000 correct (74.30)

Maximum accuracy attained: 77.0  
Checking accuracy on test set  
Got 7488 / 10000 correct (74.88)

[58]: 74.88



## 2.14 Resnet with Attention (5 points)

```
[92]: ## Resnet with Attention

learning_rate = 1e-3

# TODO: Use the above Attention module after the 2nd resnet block i.e. after
↪self.layer2.
class ResNetAttention(nn.Module):

    def __init__(self, block, layers, img_channels=3, num_classes=100,
    ↪batchnorm=False):
        super(ResNetAttention, self).__init__() #layers = [1, 1, 1, 1]
        self.in_channels = 64
        self.conv1 = nn.Conv2d(img_channels, 64, kernel_size=7, stride=2,
    ↪padding=3)
        self.bn1 = nn.BatchNorm2d(64)
        self.relu = nn.ReLU()
        self.maxpool = nn.MaxPool2d(kernel_size=3, stride=2, padding=1)
        self.batchnorm = batchnorm
        self.layer1 = self.make_layer(block, layers[0], out_channels=64,
    ↪stride=1, batchnorm=batchnorm)
        self.layer2 = self.make_layer(block, layers[1], out_channels=128,
    ↪stride=1, batchnorm=batchnorm)
        self.attention = Attention(128)
        self.layer3 = self.make_layer(block, layers[2], out_channels=256,
    ↪stride=1, batchnorm=batchnorm)
        self.layer4 = self.make_layer(block, layers[3], out_channels=512,
    ↪stride=2, batchnorm=batchnorm)

        self.averagepool = nn.AdaptiveAvgPool2d((1, 1))
        self.fc = nn.Linear(512, num_classes)

    def forward(self, x):

        x = self.conv1(x)
        if self.batchnorm:
            x = self.bn1(x)
        x = self.relu(x)
        x = self.maxpool(x)
        x = self.layer1(x)
        x = self.layer2(x)
        x = self.attention(x)
        x = self.layer3(x)
        x = self.layer4(x)
        x = self.averagepool(x)
```

```

        x = x.reshape(x.shape[0], -1)
        x = x.reshape(x.shape[0], -1)
        x = self.fc(x)

    return x

    def make_layer(self, block, num_blocks, out_channels, stride,
↳batchnorm=False):
        downsampler = None
        layers = []
        if stride != 1 or self.in_channels != out_channels:
            downsampler = nn.Sequential(nn.Conv2d(self.in_channels,
↳out_channels, kernel_size = 1, stride = stride), nn.
↳BatchNorm2d(out_channels))

        layers.append(block(self.in_channels, out_channels, downsampler,
↳stride, batchnorm=batchnorm))

        self.in_channels = out_channels

        for i in range(num_blocks - 1):
            layers.append(block(self.in_channels, out_channels))

        return nn.Sequential(*layers)

class ResNetAttention10block(nn.Module):

    def __init__(self, in_channels, out_channels, downsampler = None, stride =
↳1, batchnorm=False):

        super(ResNetAttention10block, self).__init__()
        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size = 3,
↳padding = 2)
        self.bn1 = nn.BatchNorm2d(out_channels)
        self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size = 3,
↳stride = stride)
        self.bn2 = nn.BatchNorm2d(out_channels)
        self.downsampler = downsampler
        self.relu = nn.ReLU()
        self.batchnorm = batchnorm

    def forward(self, x):

        residual = x

```

```

        x = self.conv1(x)
        if self.batchnorm:
            x = self.bn1(x)
        x = self.relu(x)
        x = self.conv2(x)
        if self.batchnorm:
            x = self.bn2(x)
        x = self.relu(x)

        if self.downsampler:
            residual = self.downsampler(residual)

        return self.relu(residual + x)

def ResNetAttention10(num_classes = 100, batchnorm= False):

    return ResNetAttention(ResNetAttention10block, [1, 1, 1, 1],
    ↪ num_classes=num_classes, batchnorm=batchnorm)

model = ResNetAttention10()

optimizer = optim.Adam(model.parameters(), lr=learning_rate)

train(model, optimizer, epochs=10)

```

Epoch 0, Iteration 0, loss = 4.6545  
 Checking accuracy on validation set  
 Got 98 / 1000 correct (9.80)

Epoch 0, Iteration 100, loss = 1.7813  
 Checking accuracy on validation set  
 Got 301 / 1000 correct (30.10)

Epoch 0, Iteration 200, loss = 1.5791  
 Checking accuracy on validation set  
 Got 466 / 1000 correct (46.60)

Epoch 0, Iteration 300, loss = 1.3994  
 Checking accuracy on validation set  
 Got 482 / 1000 correct (48.20)

Epoch 0, Iteration 400, loss = 1.3925  
 Checking accuracy on validation set  
 Got 496 / 1000 correct (49.60)

Epoch 0, Iteration 500, loss = 1.2355  
 Checking accuracy on validation set

Got 516 / 1000 correct (51.60)

Epoch 0, Iteration 600, loss = 1.5204  
Checking accuracy on validation set  
Got 507 / 1000 correct (50.70)

Epoch 0, Iteration 700, loss = 1.3769  
Checking accuracy on validation set  
Got 584 / 1000 correct (58.40)

Epoch 1, Iteration 0, loss = 0.8449  
Checking accuracy on validation set  
Got 538 / 1000 correct (53.80)

Epoch 1, Iteration 100, loss = 1.1106  
Checking accuracy on validation set  
Got 593 / 1000 correct (59.30)

Epoch 1, Iteration 200, loss = 1.0164  
Checking accuracy on validation set  
Got 590 / 1000 correct (59.00)

Epoch 1, Iteration 300, loss = 1.0507  
Checking accuracy on validation set  
Got 581 / 1000 correct (58.10)

Epoch 1, Iteration 400, loss = 1.0679  
Checking accuracy on validation set  
Got 615 / 1000 correct (61.50)

Epoch 1, Iteration 500, loss = 1.1800  
Checking accuracy on validation set  
Got 641 / 1000 correct (64.10)

Epoch 1, Iteration 600, loss = 1.2314  
Checking accuracy on validation set  
Got 656 / 1000 correct (65.60)

Epoch 1, Iteration 700, loss = 0.9316  
Checking accuracy on validation set  
Got 683 / 1000 correct (68.30)

Epoch 2, Iteration 0, loss = 0.7903  
Checking accuracy on validation set  
Got 647 / 1000 correct (64.70)

Epoch 2, Iteration 100, loss = 1.0066  
Checking accuracy on validation set

Got 663 / 1000 correct (66.30)

Epoch 2, Iteration 200, loss = 0.8983  
Checking accuracy on validation set  
Got 652 / 1000 correct (65.20)

Epoch 2, Iteration 300, loss = 0.8269  
Checking accuracy on validation set  
Got 653 / 1000 correct (65.30)

Epoch 2, Iteration 400, loss = 0.7509  
Checking accuracy on validation set  
Got 665 / 1000 correct (66.50)

Epoch 2, Iteration 500, loss = 0.9236  
Checking accuracy on validation set  
Got 653 / 1000 correct (65.30)

Epoch 2, Iteration 600, loss = 0.8746  
Checking accuracy on validation set  
Got 670 / 1000 correct (67.00)

Epoch 2, Iteration 700, loss = 0.6206  
Checking accuracy on validation set  
Got 702 / 1000 correct (70.20)

Epoch 3, Iteration 0, loss = 0.6253  
Checking accuracy on validation set  
Got 698 / 1000 correct (69.80)

Epoch 3, Iteration 100, loss = 0.7086  
Checking accuracy on validation set  
Got 731 / 1000 correct (73.10)

Epoch 3, Iteration 200, loss = 0.7628  
Checking accuracy on validation set  
Got 714 / 1000 correct (71.40)

Epoch 3, Iteration 300, loss = 0.5393  
Checking accuracy on validation set  
Got 720 / 1000 correct (72.00)

Epoch 3, Iteration 400, loss = 0.7262  
Checking accuracy on validation set  
Got 713 / 1000 correct (71.30)

Epoch 3, Iteration 500, loss = 0.8739  
Checking accuracy on validation set

Got 723 / 1000 correct (72.30)

Epoch 3, Iteration 600, loss = 0.7431  
Checking accuracy on validation set  
Got 726 / 1000 correct (72.60)

Epoch 3, Iteration 700, loss = 0.7139  
Checking accuracy on validation set  
Got 708 / 1000 correct (70.80)

Epoch 4, Iteration 0, loss = 0.4678  
Checking accuracy on validation set  
Got 729 / 1000 correct (72.90)

Epoch 4, Iteration 100, loss = 0.4696  
Checking accuracy on validation set  
Got 739 / 1000 correct (73.90)

Epoch 4, Iteration 200, loss = 0.4469  
Checking accuracy on validation set  
Got 740 / 1000 correct (74.00)

Epoch 4, Iteration 300, loss = 0.4950  
Checking accuracy on validation set  
Got 723 / 1000 correct (72.30)

Epoch 4, Iteration 400, loss = 0.7305  
Checking accuracy on validation set  
Got 714 / 1000 correct (71.40)

Epoch 4, Iteration 500, loss = 0.4393  
Checking accuracy on validation set  
Got 729 / 1000 correct (72.90)

Epoch 4, Iteration 600, loss = 0.6803  
Checking accuracy on validation set  
Got 747 / 1000 correct (74.70)

Epoch 4, Iteration 700, loss = 0.7277  
Checking accuracy on validation set  
Got 739 / 1000 correct (73.90)

Epoch 5, Iteration 0, loss = 0.3798  
Checking accuracy on validation set  
Got 743 / 1000 correct (74.30)

Epoch 5, Iteration 100, loss = 0.6093  
Checking accuracy on validation set

Got 735 / 1000 correct (73.50)

Epoch 5, Iteration 200, loss = 0.4920  
Checking accuracy on validation set  
Got 709 / 1000 correct (70.90)

Epoch 5, Iteration 300, loss = 0.6187  
Checking accuracy on validation set  
Got 748 / 1000 correct (74.80)

Epoch 5, Iteration 400, loss = 0.4713  
Checking accuracy on validation set  
Got 765 / 1000 correct (76.50)

Epoch 5, Iteration 500, loss = 0.6515  
Checking accuracy on validation set  
Got 768 / 1000 correct (76.80)

Epoch 5, Iteration 600, loss = 0.6915  
Checking accuracy on validation set  
Got 765 / 1000 correct (76.50)

Epoch 5, Iteration 700, loss = 0.5748  
Checking accuracy on validation set  
Got 756 / 1000 correct (75.60)

Epoch 6, Iteration 0, loss = 0.5636  
Checking accuracy on validation set  
Got 762 / 1000 correct (76.20)

Epoch 6, Iteration 100, loss = 0.4709  
Checking accuracy on validation set  
Got 760 / 1000 correct (76.00)

Epoch 6, Iteration 200, loss = 0.2997  
Checking accuracy on validation set  
Got 774 / 1000 correct (77.40)

Epoch 6, Iteration 300, loss = 0.3733  
Checking accuracy on validation set  
Got 773 / 1000 correct (77.30)

Epoch 6, Iteration 400, loss = 0.4140  
Checking accuracy on validation set  
Got 738 / 1000 correct (73.80)

Epoch 6, Iteration 500, loss = 0.7153  
Checking accuracy on validation set

Got 751 / 1000 correct (75.10)

Epoch 6, Iteration 600, loss = 0.5156  
Checking accuracy on validation set  
Got 768 / 1000 correct (76.80)

Epoch 6, Iteration 700, loss = 0.4074  
Checking accuracy on validation set  
Got 767 / 1000 correct (76.70)

Epoch 7, Iteration 0, loss = 0.2979  
Checking accuracy on validation set  
Got 777 / 1000 correct (77.70)

Epoch 7, Iteration 100, loss = 0.3108  
Checking accuracy on validation set  
Got 779 / 1000 correct (77.90)

Epoch 7, Iteration 200, loss = 0.3539  
Checking accuracy on validation set  
Got 769 / 1000 correct (76.90)

Epoch 7, Iteration 300, loss = 0.5528  
Checking accuracy on validation set  
Got 760 / 1000 correct (76.00)

Epoch 7, Iteration 400, loss = 0.4538  
Checking accuracy on validation set  
Got 779 / 1000 correct (77.90)

Epoch 7, Iteration 500, loss = 0.6501  
Checking accuracy on validation set  
Got 759 / 1000 correct (75.90)

Epoch 7, Iteration 600, loss = 0.2808  
Checking accuracy on validation set  
Got 785 / 1000 correct (78.50)

Epoch 7, Iteration 700, loss = 0.4682  
Checking accuracy on validation set  
Got 755 / 1000 correct (75.50)

Epoch 8, Iteration 0, loss = 0.1710  
Checking accuracy on validation set  
Got 772 / 1000 correct (77.20)

Epoch 8, Iteration 100, loss = 0.4158  
Checking accuracy on validation set



Got 777 / 1000 correct (77.70)

Epoch 8, Iteration 200, loss = 0.2259  
Checking accuracy on validation set  
Got 784 / 1000 correct (78.40)

Epoch 8, Iteration 300, loss = 0.4438  
Checking accuracy on validation set  
Got 784 / 1000 correct (78.40)

Epoch 8, Iteration 400, loss = 0.4175  
Checking accuracy on validation set  
Got 776 / 1000 correct (77.60)

Epoch 8, Iteration 500, loss = 0.2937  
Checking accuracy on validation set  
Got 770 / 1000 correct (77.00)

Epoch 8, Iteration 600, loss = 0.3422  
Checking accuracy on validation set  
Got 754 / 1000 correct (75.40)

Epoch 8, Iteration 700, loss = 0.2597  
Checking accuracy on validation set  
Got 780 / 1000 correct (78.00)

Epoch 9, Iteration 0, loss = 0.4149  
Checking accuracy on validation set  
Got 776 / 1000 correct (77.60)

Epoch 9, Iteration 100, loss = 0.3305  
Checking accuracy on validation set  
Got 769 / 1000 correct (76.90)

Epoch 9, Iteration 200, loss = 0.3941  
Checking accuracy on validation set  
Got 753 / 1000 correct (75.30)

Epoch 9, Iteration 300, loss = 0.2413  
Checking accuracy on validation set  
Got 761 / 1000 correct (76.10)

Epoch 9, Iteration 400, loss = 0.2372  
Checking accuracy on validation set  
Got 778 / 1000 correct (77.80)

Epoch 9, Iteration 500, loss = 0.4171  
Checking accuracy on validation set

Got 761 / 1000 correct (76.10)

Epoch 9, Iteration 600, loss = 0.4070

Checking accuracy on validation set

Got 780 / 1000 correct (78.00)

Epoch 9, Iteration 700, loss = 0.2234

Checking accuracy on validation set

Got 786 / 1000 correct (78.60)

Maximum accuracy attained: 78.60000000000001

## 2.15 Test set – run this only once

Now we test our model on the test set . Think about how this compares to your validation set accuracy.

```
[93]: AttentionResnet = model
      check_accuracy(loader_test, AttentionResnet)
```

Checking accuracy on test set

Got 7701 / 10000 correct (77.01)

[93]: 77.01

## 2.16 Inline Question 3: Rank the above models based on their performance on test dataset (15 points)

( You are encouraged to run each of the experiments (training) at least 3 times to get an average estimate )

Report the test accuracies alongside the model names. For example, 1. Vanilla CNN (57.45%, 57.99%).. etc

1. Resnet with attention (75.98%, 76.34%, 77.01%)
2. Vanilla Resnet (74.88%, 74.43%, 73.52%)
3. CNN with early attention (60.60%, 62.32%, 61.12%)
4. CNN with double attention blocks (58.31%, 60.24%, 58.62%)
5. Vanilla CNN (59.51%, 56.98%, 59.52%)
6. CNN with late attention (56.37%, 57.68%, 58.30%)

### 2.16.1 Bonus Question (Ungraded): Can you give a possible explanation that supports the rankings?

Your Answer:

Sometimes Resnet with attention performs worse than Resnet vanilla. This is because the sampling efficiency decreases as we use Resnet with attention instead of Resnet, making the network require more epochs to train. After increasing the number of training epochs, the performance of Resnet with attention stably gets better than the vanilla one.